

Peer to Peer Systems and Security

Documentation for the VoidPhone Project

Team “The Hashies”

## 1. Latest version of our architecture / structure as introduced in the midterm report

Our DHT-module is now split up into three main parts:

- dht\_server.cpp
- routing.cpp
- ssl.cpp

In the following, we explain the dht\_server.cpp file. This file is compiled to the final executable. It contains major server logic and is dependent on the other two source files. Important remark / disclaimer for grading: The ssl.cpp was almost exclusively auto generated with LLMs like ChatGPT a. o. We do not want to take credit for this file. Nonetheless, the integration, correct usage, and rewriting of several methods by hand easily summed up to 30 hrs. of effort for Jörn.

In the main function of dht\_server, we use boost to parse command line arguments. We pass the following arguments to start a server.

### Command line (Input) Arguments:

- host-address: The IP address by which the host will be reachable for peers. This can be an IPv4 or IPv6 address.
- module-port: The port on which we listen for API request from other modules (DHT PUT, DHT GET)
- p2p-port: The port on which we listen for RtiPCs by other peers.

Only passing these arguments will start the DHT server as a new network. To connect to an existing network, we need to pass two additional arguments to reach a node in the existing network:

- peer-address: IP of one peer to join their network (again, IPv4 or IPv6)
- peer-port: Corresponding p2p-port of the peer

If we know multiple addresses of potential nodes, we can try to start the server with any node (for this, several executions with modified command line arguments may be needed). As soon as one is responsive, the server joins the network.

Additional flags such as `-l` for logging level and `-v` for verification of certification can be read from the help page.

We can also pass `--help` to see examples of usage and description for the available arguments.

### Threading scheme:

Our server currently uses two threads; One main thread does all logical decisions, all state-keeping of connections and answering requests and even implements an asynchronous task queue, in which we can emplace work needed to be done in the future. One additional background thread is used that periodically purges the `local_storage` of our node, which contains the data previously DHT\_PUT by users. It is used to comply with the `TimeToLive` requirement, which each DHT\_PUT contains and was specified. The use of this extra thread was a convenience-driven decision, as the `local_storage` is rarely accessed for most work like routing and session management. Using a mutex that guarantees resource exclusivity during access times to `local_storage` suffices and does not create any noticeable delays.

### Routing, Routing Table, Kademlia:

Next, we are generating a `RoutingTable` object, which will handle all our needs related to our `K_Buckets` and node handling. This will also randomly generate our node ID. More on that later, same as `epoll`, which we set up following that.

After the `RoutingTable`, if a peer was passed, we try to connect to its network. To do that, we are sending a "find\_node" RPC to the given peer with our own network. The peer will then look in its own network to find peers close to our node ID and return them to us with the `Node` struct, made up of IP, Port and node ID of the peer. We will receive these node IDs with an `epoll_wait` call. Now, to expand our network further than up to 20 nodes which we got from the peer, we are then doing a so called "bucket refresh": For each of our buckets we look for random keys in the range of each bucket. We do this until we don't add any more nodes than before. In a very large network, this could mean that we filled all of our buckets, but for our purposes we usually find the few peers we started, for example with a network of about 30 peers we should get the contact details and node ID of every peer in the network.

After gathering the `Node` struct of our peers, our main loop starts. We have set up two ports, one for the module API and one for p2p traffic. For both, we create one socket that is listened on with `epoll`. Every time we receive a new request on either socket, an `epoll` event is triggered.

When a new node is contacting us, it's on the module API socket or on the P2P socket. To distinguish between them we pass a "ConnectionType" to the function that set the new connection up. The `ConnectionType` for "accept\_new\_connection" is either `MODULE_API` or `P2P`.

`accept_new_connection` accepts and builds up the TCP connection on a new socket, sets up an SSL tunnel and adds a new entry to the most important structure of our networking: the `connection_map`. More on TLS/SSL later. The map saves a `ConnectionInfo` object for every socket we create. The `ConnectionInfo` Object's main purpose is therefore session management. It contains the `ConnectionType`, client address and port, RPC ID (explained later), TLS data, and two buffers: one for incoming and one for outgoing traffic.

When we now get a new request on that newly registered socket, another EPOLL event is triggered. We then read the incoming traffic and pass it to the `try_processing` function, which parses the header. We then look whether the traffic came in on a module API socket or a p2p socket (indicated by the `ConnectionType` in the `ConnectionInfo`) and check the validity of the header.

Now comes the main part of the networking: For every request type we have a handle and send function. These are sending and receiving nonblockingly, so in a way that enables other traffic to flow even when we have larger requests. This is implemented by never calling `read()` or `write()` on sockets directly, even though marked nonblocking.

Epoll-in events flush the kernel-side socket buffer into our session-management's receive-buffer. Handling of the content is postponed to the next Epoll-in event if the request is not fully present in our receive buffer (handle-functions detect this incompleteness of messages). Epoll-out events flush our session-management's send-buffer into the kernel-side socket buffer. If we can't flush everything, we flush only partially and leave the rest for the next Epoll-out event (the forged message can fully or partially reside in the send-buffer at any time). The flushing operations are the only ones which call `read()` or `write()` on the sockets. All our defined functions only directly interact with the session-management buffers, never with the kernel-side socket buffers.

If we are receiving a request on a socket of `ConnectionType` Module API, we continue with the `parse_API_request` function, otherwise with the `parse_P2P_request` function.

For the former, we decide between a PUT request and GET request, for the latter, we have 9 RPC functions: PING, STORE, FIND\_NODE, FIND\_VALUE, a reply call for each and ERROR, which tells the sender of an RPC if the call failed because of a malformed request (e.g. forgot key with a FIND\_NODE request) or an internal server error.

Every RPC sends the normal header (4 bytes total: size and type) followed by the RPC header (66 bytes total: node ID, port, and RPC ID). The RPC ID is randomly generated ID that should prevent address forgery, which is tested for when the call is received. This is a change to the protocol outlined in the midterm report, as we initially forgot to pass the node ID and port (the port because we send requests from a new socket instead of the p2p port). The security implication here is that in production, where malicious nodes could join, we assume that a stranger's peer's nodes are started on a different IP address. For testing, we start docker containers in a docker network. There, every peer has its own IP address, which solves this problem.

For every RPC we have 4 functions (coming one after another in the code):

- `forge_DHT_RPC_<method>`

Generate an RPC ID and fill in the appropriate fields.

- `handle_DHT_RPC_<method>`

Read the RPC ID and save it into the `connection_map`, also read node ID, Port and IP. Test if we already know the peer. If not, we add him to our `RoutingTable`.

- `forge_DHT_RPC_<method>_reply`

We look up the previously received RPC ID and fill in the fields for replying.

- `handle_DHT_RPC_<method>_reply`

We first check whether the RPC ID was the one which was previously sent, then we process the response.

In each of these functions we either use the `build_DHT_header`, `build_RPC_header` and `write_body` functions (in the forge functions) or the `read_body` (in the handle functions). The `read_body` and `write_body` functions move data between the message buffer and the given data structures and are called for every field in the message (key, value etc.). For the message itself, we use a vector of unsigned chars, abstracted away in the "Message"-type. Unsigned, because they display a more appropriate decimal number when debugging and are therefore easier to handle during development.

To respond to requests, we have the second most important map in our server: The `request_map`.

Here, we save a shared pointer to a request object, that is shared among all participants of the request. For example, if we do a network expansion, every socket that is part of that operation is put into the `request_map.at(socketfd)` and gets passed a shared pointer to the larger "Request" struct, that saves the current state of the entire operation. Every time a socket gets a node reply, it looks into the `request_map` and then checks in the linked Request object what it should do next based on the state inside. This way, we do not need any multi-threading, because we have turned our entire server in a complete state machine, which worked much better than what was previously implemented. (See the last version of this document in git if you are interested. We can not recommend reading that though...). This stateful rewrite of the server caused some problems because of lacking time though (see known issues). We still are convinced that this rewrite was a good idea, since for us it was very important to value appropriate coding style and clean logic.

## Local storage and retrieval of values

Locally, each Node stores the values for which it is responsible in a simple hash map. This map uses the hashes of the values maps it to the value to store. The hashes are in the 256-bit space, as well as the Kademlia-IDs, so we store the values only on the `k` closest (->XOR-metric) Nodes, where `k` is the replication factor. The specification also required the Nodes to respect the time to live of a value in the storage that is provided in the TTL field of the PUT request. For our implementation, this is done via an extra thread that periodically purges the storage and all its outdated values. The thread is not actively waiting, but elegantly waits on a condition variable to not consume resources. The condition variable was necessary as we needed to be able to join the purging thread at any point of time in our application. This is essential, as the user can send an interrupt or kill signal to the server at any time, so a signal handler was setup to perform trailing clean up logic. As this logic also joins the purging thread, we use a condition variable to notify the purging thread that it's time to return has come.

## How does the server work?

After this DHT-related introduction to our structure, we want to introduce you to some design decisions regarding the pure server functionality of our DHT-module.

One small introduction to this was the previously mentioned session management via `connection_map` and its contained `ConnectionInfo` objects, each one representing a singular session. Let us start with the big picture: TLS, a custom handshake protocol enhancing security and `epoll`:

### How are sessions initiated?

As our server implements TLS, we need a TLS client and a TLS server role in each connection, even though we treat everyone equally in the Peer-to-Peer network setting. For clarity: A client is typically a user or other server which has a request for a serving server. The client `connect()`s to the server. The server is the peer in the connection which actively `listen()`s for new connections and `accept()`s them on a different port.

For implementing TLS, a slightly custom session instantiation protocol was used. On the start of a server, the server generates an SSL certificate which contains the server IP and the node ID from Kademlia. Both values are constant for the lifetime of the server. Therefore, the certificate can be generated before sessions are accepted, especially since the certificate will be provided on every new session, this keeps computational effort and memory usage negligible.

Maybe a question arises, namely, how does one generate a certificate for oneself? Principles like “root of trust” hinder. After reflecting for quite some time, for our P2P-Network setting, a centralized trusted instance like a Certification Authority (CA) or ticket server were out of question. PKI’s as used for e.g. the internet are also hard to realize, as for this to work, certain peers have substantially more power (trust) than others. We wanted to keep the network a network of equals. Web-of-Trust based designs like gathering trust from known peers were also thought of, but hard to realize. Finally, we stuck to X509 certificates but decided on self-signing them. This, of course, does not impact the trustworthiness of the certificates in any way, but with the “Trust On First Use” (TOFU) principle, we can initially assume the credibility of the certificates and their proof of authenticity of their creators. More on this certificate usage below.

To dive deeper into the session instantiation protocol, we consider three distinguishable parts of the protocol: TCP Handshake, Server Provides Certificate to Client, TLS Handshake.

#### 1.1. TCP Handshake:

Typical socket-layer TCP handshake: Server and client create sockets with `socket()`, the server `bind()`s and `listen()`s on dedicated port, the client `connect()`s to the dedicated port and the server finally `accept()`s the connection (and answers via a different port than his dedicated `listen()` port). This entire TCP handshake is also performed fully non-blocking by keeping additional state and waiting on `epoll` for the 1.5 round-trip to successfully complete after initiating it on the client-side.

#### 1.2. Server Provides Certificate to Client:

After the TCP-Handshake succeeded, the first custom operation is initialized. The server provides its pre-generated, self-signed certificate which contains two major X509 v.3 custom extensions: The IPv6 field and the Kademlia-ID in a field. The port is not contained (the three would make a true identifier of any server), as the self-signing effort for each new port would be infeasible. Before transmission, the certificate is preceded by 4 bytes of unsigned 32-bit integer in network byte order which indicate the length of the succeeding certificate. A typical length-prefixed send, as certificates could vary arbitrarily in length. During debugging, we realized that almost all our certificates are 1245 bytes long.

As soon as the full certificate is present on the client's side, the client does the heavy lifting part. It takes the certificate and compares it against all previously received certificates (which it bookkeeps for its entire lifetime). If the Kademlia-ID is recognized (contained in another certificate), the client gets suspicious. Theoretically, the provided certificate could be valid if the old, stored certificate's subject (and issuer, as self-signed) churned out of the network. Then, another peer with a different IP could replace the churned peer under the same Kademlia-ID.

Due to this special case, the client performs an RPC\_ping to the owner of the old server certificate. If the old owner is still reachable, the new connection is aborted, as the new certificate could be maliciously generated by choosing a Kademlia-ID instead of randomly generating it, and therefore trying to spoof the identity of a possibly long-term established P2P-network server. If the old owner is unreachable, the old certificate is considered deprecated and new "Trust On First Use" (TOFU) can be established in the freshly received certificate. Finally, after the certificate was sent by the server and read (& verified) by the client, both parties can start with the typical SSL-Handshake.

### 1.3. Typical TLS Handshake:

The last part of our custom protocol is a typical TLS handshake. As we use OpenSSL, this is done with SSL\_accept() and SSL\_connect(). The TLS version our server uses is TLS 1.3, state of the art. We use 2048-bit private key-length and RSA key-generation. Additionally, as a necessity of our nonblocking sending and receiving, even this TLS handshake is implemented nonblocking via epoll event notification and handle\_custom\_ssl\_protocol() for automaton-based SSLState transition. Once CONNECTED or ACCEPTED, depending on whether client or server succeeded in the handshake, we automatically flush out all future send\_buffer contents via nonblocking SSL\_write and flush all future incoming data in our recv\_buffer via nonblocking SSL\_read. For data being written to or from our ConnectionInfo Struct (emplaced in the connection\_info map), the use of SSL encryption or not is fully transparent. This implementation design improved development drastically.

In the folder "tests" we have storage and routing unit tests

## **2. Software Documentation**

- Building / running dependencies

1. If you want to run the program on your local machine, the following steps are necessary. For docker-usage, refer to 2.

To build, we use CMAKE, make sure you have a version > 3.13.

Apart from that, we need a compiler for C++, e.g. g++ or clang, that can compile C++23.

Our dependencies are Catch2, Boost, OpenSSL and sdlog, of which only Catch2 is installed automatically with CMAKE. The other two can be installed via your package manager, e.g. on ubuntu with:

```
sudo apt-get install libboost-all-dev
```

```
sudo apt-get install libssl-dev
```

```
sudo apt-get install libspdlog-dev
```

First, we clone the repository:

```
https://gitlab.lrz.de/netintum/teaching/p2psec_projects_2024/DHT-16
```

We then need to create a build folder:

```
mkdir build && cd ./build
```

configure cmake in the build folder:

```
cmake ..
```

and build:

```
cmake --build . --target dht_server
```

This should download catch2, then build the project and generate a "dht\_server" executable.

To run, we first need to create a network on at least one node. Let's assume we create a network with IPv4 local addresses for now, one peer on 192.168.0.42 and one peer on 192.168.0.43. The arguments -a -m and -p are the host address, module API port and P2P port of our new node, respectively. Therefore, to start a new network on a node with IP 192.168.0.42, call in the build folder:

```
./dht_server -a 192.168.0.42 -m 7401 -p 7402
```

Now, to join the network on a node with IP 192.168.0.43, we pass the IP and p2p-port 7402 of the just started peer with -A and -P, respectively.

```
./dht_server -a 192.168.0.43 -m 7401 -p 7402 -A 192.168.0.42 -P 7402
```

To test the server on your local machine you can create test nodes with:

```
./dht_server -a ::1 -m 7401 -p 7402
```

```
./dht_server -a ::1 -m 7403 -p 7404 -A ::1 -P 7402
```

```
./dht_server -a ::1 -m 7405 -p 7406 -A ::1 -P 7402
```

```
./dht_server -a ::1 -m 7407 -p 7408 -A ::1 -P 7402
```

For further testing, we also created the folder tests/integration. This tests file tests/integration/test\_multiple\_peers.py that can be used to start and communicate between several nodes.

## 2. Docker-based testing /starting the DHT module(s):

If you want to use Docker for testing, please view the Dockerfile in our projects root directory. It uses Ubuntu 22.04 as a base image for avoiding the "works-on-my-machine" effect. With it, you can start peers themselves in separate containers. First, view our testing recommendation: Use our python script tests/integration/start\_docker\_containers.py to instantiate and run predefined tests and investigate results as well as logging. Second, not recommended but easily possible for extra testing effort, build the docker network and run the peer containers yourself by following this guide:

Setup custom docker network for the application...

```
docker network inspect dht_network
```

Is the network up? If not, run...

```
docker network create --subnet=172.20. 0.0/16 dht_network
```

Starting 172.20.0.2:< P2P port> to create a network...

```
docker run
  -p          <Expos.      ModuleAPI      port>:<Expos.      ModuleAPI      port>
  -p          <Expos.      P2P          port>:          <Expos.      P2P          port>
  -t
  --name
  --network
  --ip
  dht_swarm
  /server/build/dht_server -a 172.20.0.2 -m <Expos. ModuleAPI port> -p <Expos. P2P port> -l info
run
                                dht_swarm_0
                                dht_network
                                172.20.0.2
```

Both -p options expose the port of the container to the host machine. The -t option shows you the terminal output of the program inside the container. The --name simply sets a name for the container, --network specifies the name of the docker network that shall be used. Network specifies the network, ip the IP. dht\_swarm the image, and the rest is the application execution call with the according arguments. Refer the --help for questions about our executable.

To join a second container to the peer, view this example command:

```
docker run
  -t
  --name dht_swarm_1
  --network dht_network
```



```
--ip 172.20.0.3 dht_swarm /server/build/dht_server -a 172.20.0.3 -m <Module API port> -p  
<P2P port> -A 172.20.0.2 -P <P2P port> -l info
```

The ports of this new peer are not exposed to the host machine. You could do this to send new module API requests to a second peer but be aware of not using the same host-ports when exposing peer ports via the `-p <host - port>:<container - port>`. For any further reading regarding the use-case, view the `start_docker_containers.py` script in the `tests/integration/` folder. The print mode of this script will print example commands that can directly be copy pasted.

During manual testing, we recommend setting all Module-API ports to 7401, and all P2P-Ports to 7402 and only exposing one or two containers' ports to the host machine, as exposed ports and unexposed are easily mixed and confused. By doing this, the containers will only differ by IP in the docker network, not by ports.

- Known issues:

When TLS is set to `SSL_VERIFY_PEER` for the client of our TLS connection, sometimes an error occurs. We tried to debug it for hours of hours, this bug even motivated us to radically rewrite the entire server side on the last week before the updated submission deadline and implement detailed logging, mainly to avoid these errors where our deviation from protocol could cause fatal errors during connection setup phase.

The main challenge was writing everything non-blocking, so always waiting for callback, also maintaining excessive state for each session. Regarding the already addressed bug, we currently have only one suspicion which we think we already ruled out: The server side unintentionally send some user-data, maybe also the certificate which is sent by the server in between the TCP and SSL handshake, which is in turn interpreted as part of the SSL handshake on the client side of our session. This leads the internal SSL library to parse the certificate or the other user data as SSL-handshake related data that is otherwise expected from an `SSL_accept()` call of the server-side.

We had similar misbehavior of our protocol some time ago and fixed it, but the error back then was “error:0A00010B:SSL routines::wrong version number”. This made sense, as the Client-Hello requests the server certificate during the handshake, expects the very first bytes of the handshake-certificate to be the version number of TLS, but instead receives user-data. Currently, the client side error is “error:0A000086:SSL routines::certificate verify failed”, indicating that only the certificate validation fails. The server error is “error:0A000418:SSL routines::tlsv1 alert unknown ca “, indicating that the certification authority is unknown, even though we have the identical certificate contained in our trust-store. We wrote methods to dump the current certificate store of our client and also to dump the initially generated certificate of our server. They match, byte for byte, formatted identically. During debug-inspections, the transmitted byte count is identical, everything seems perfect. We invested over 10 hours debugging only to fix this singular bug, back then before we had rewritten the server drastically and after we rewrote it. The part that is truly hard to debug is that we cannot inspect the data that relies on the kernel-side buffer of our socket. If we flush it into our receive buffer (fetch the data from the socket into our user-data buffer), we cannot continue with the SSL protocol as this data is needed by `SSL_accept/connect` to reside on the file-descriptor's kernel-side buffer so that the SSL library can consume the data for internal decision logic. Meanwhile, if it was user data, we could not inspect it as it must reside on the kernel-side buffer of our socket...

Unfortunately, our attempts were without success. Therefore, we needed to set verification of the server certificate to `SSL_VERIFY_NONE` on the client side for most tests. We still a test to showcase this bug. For this, the `SSL_VERIFY_<PEER/NONE>` is passed as a command line option which defaults to `SSL_VERIFY_NONE`. The entire traffic itself is still TLS encrypted; we only cannot ensure that the provided certificates are signed correctly / contained in our application's X509 store and therefore trusted.

The second issue is that our logger normally prints colorful to the console, but when running the docker container and redirecting the console output via `-t`, the logging looks like `"[[32minfo[m]"` instead of a grey color that says `[info]`. There is currently no workaround for this. If the program is executed directly, without docker abstraction, the colorful fonts work.

Another big issue is sadly, that through a big rewrite towards the end we introduced some bugs in the main kademia logic that would have taken more time to fix. On the other side, this big rewrite made the logic much more robust. We can currently start 100 docker containers without a problem and they all connect to the network, which was not possible before the rewrite. Some nodes just blocked on a connect call before or through weird multithreading calls just stopped responding. Now, with the main logic being completely single threaded, we have avoided an entire category of problems. All of this came at a great cost though: Even though we have invested an entire week each, without doing much other than coding, we could not fix all new bugs. We were both very saddened by this, as it was the most time we've ever spent on a on a coding project und would have loved to bring the kademia logic to an even better working condition, but that is the "curse" of deadlines.

### 3. Future Work

In the future, the `epoll`-event-loop based server can be multithreaded. Achieving this should be realistic, as `epoll` provides really good threading-support, such that individual events can be worked on by a worker-thread pool. Until now, we have not sided on implementing multithreading as this makes debugging even more difficult and currently, we still have the abovementioned bug of SSL certificate validation which we could not fix yet.

An additional goal would be to detect malicious nodes in the network. During our concurrent lookup, if a singular node repeatedly replies with DHT values which differ from the majority of answers, we can grow suspicion on these nodes and remove them if the suspicion is above a certain threshold. Also, if a peer replies repeatedly with syntax-errored reply data-frames, we could blacklist this node as it might seem to send random traffic just to exhaust bandwidth and computation power. These are future considerations whose implementation would make the nodes, and thereby the network more robust.

Lastly, a nice future work would be to fix the bug regarding the SSL certificate validation, as the SSL certificate generation and verification works in some cases but seems to fail arbitrarily in others. We still provided a test in the python script that enforces one-sided certificate verification of the server-certificate on the client side. Fixing this bug is a future goal for employing authentication during messaging. The general message passing logic of the server (write on user buffer->flush to socket, received on socket buffer-> flush to user buffer) works flawlessly, excluding the already mentioned SSL certificate verification step.

### 4. Workload distribution — Who did what

In the initial workload separation Marius planned to handle the RoutingTable part, while Jörn wanted to focus on the networking in the dht\_server. After the main server implementation with epoll and the epoll event loop was finished by Jörn, which was a lot of effort, Marius focused on implementing K Buckets and the Bucket List that contains the K Buckets. Later, Jörn decided to deal with the implementation of all security related aspects, including creating and self-signing certificates, initiating TLS connections, and sending and receiving secure messages, while Marius implemented the rest of the server, e.g. handle and forge functions, the RPC's and crawl logic. In the final week of the project, Jörn contributed to major rewrites of the server structure while Marius completely rewrote the processing of requests. Both tasks were very time-intensive and can be compared in workload intensity.

## **5. Effort spent for the project (individually)**

Jörn's main contributions, the networking logic and TLS were both an immense effort as there was no clear-cut guide on how this was to be implemented. Apart from that, we were both not very familiar with signing certificates and implementing TLS, so he used the internet for help where needed and spent a lot of time iterating over the functionality. This resulted in about 1700 additional lines of code just for implementing the TLS logic (ssl.cpp file, 800 of the 1700 LoC, for which we want no credit, as it is, as previously mentioned, LLM generated), which took him in total about 65 hours. Then, during the last week before the submission, Jörn dedicated his life into rewriting the server functionality entirely, making structural changes such as propagating errors up instead of handling them deeply nested, simplifying the epoll-structure, adding more state for connection-status bookkeeping, setting up docker, and implementing extensive logging with sdlog for the entire project, setting up an async task queue etc. This easily summed up to 800+ LoC, and amounted to at least 60 hours of work time, averaging more than 8 hours work per day. This is only an estimate, we do not want to exaggerate, please view the commit history. Adding to this came about 15 hours of work on other things.

Meanwhile, Marius' part was also time-consuming but initially more predictable, as it was relatively clear from the specification and the kademia paper what needed to be implemented. His initial work on the routing and kademia protocol took about 55 hours. The rewrite of the entire server that was mentioned earlier was done mainly within the last 5 days, where he worked over 12 hours per day on average, sometimes working till the morning and sleeping little, until the features were finished. A lot of the past work was completely redone so the main server logic can be completely single threaded and still operate concurrently. The rewrite was still necessary, as it greatly improved our code.

Considering all of the work over the last ~2 months, Jörn had a time effort on average of about 17.5 hours per week while Marius an effort of about 13 hours per week. Both of us did vastly more work towards the end of the working period, again, view the commit history.